

# Information Systems

## Chapter 5:

## Query Processing

Rita Schindler | TU Ilmenau, Germany

[www.tu-ilmenau.de/dbis](http://www.tu-ilmenau.de/dbis)

# Query Processing

---

- ▶ Basic Steps
- ▶ Cost Model
- ▶ Join Operations
- ▶ Other Operations
- ▶ Evaluation of Expressions

# Query Processing and Optimization

---

## Task

- ▶ Find an optimal (efficient) evaluation plan for a descriptive (SQL) query without accessing data files (except data dictionary tables)

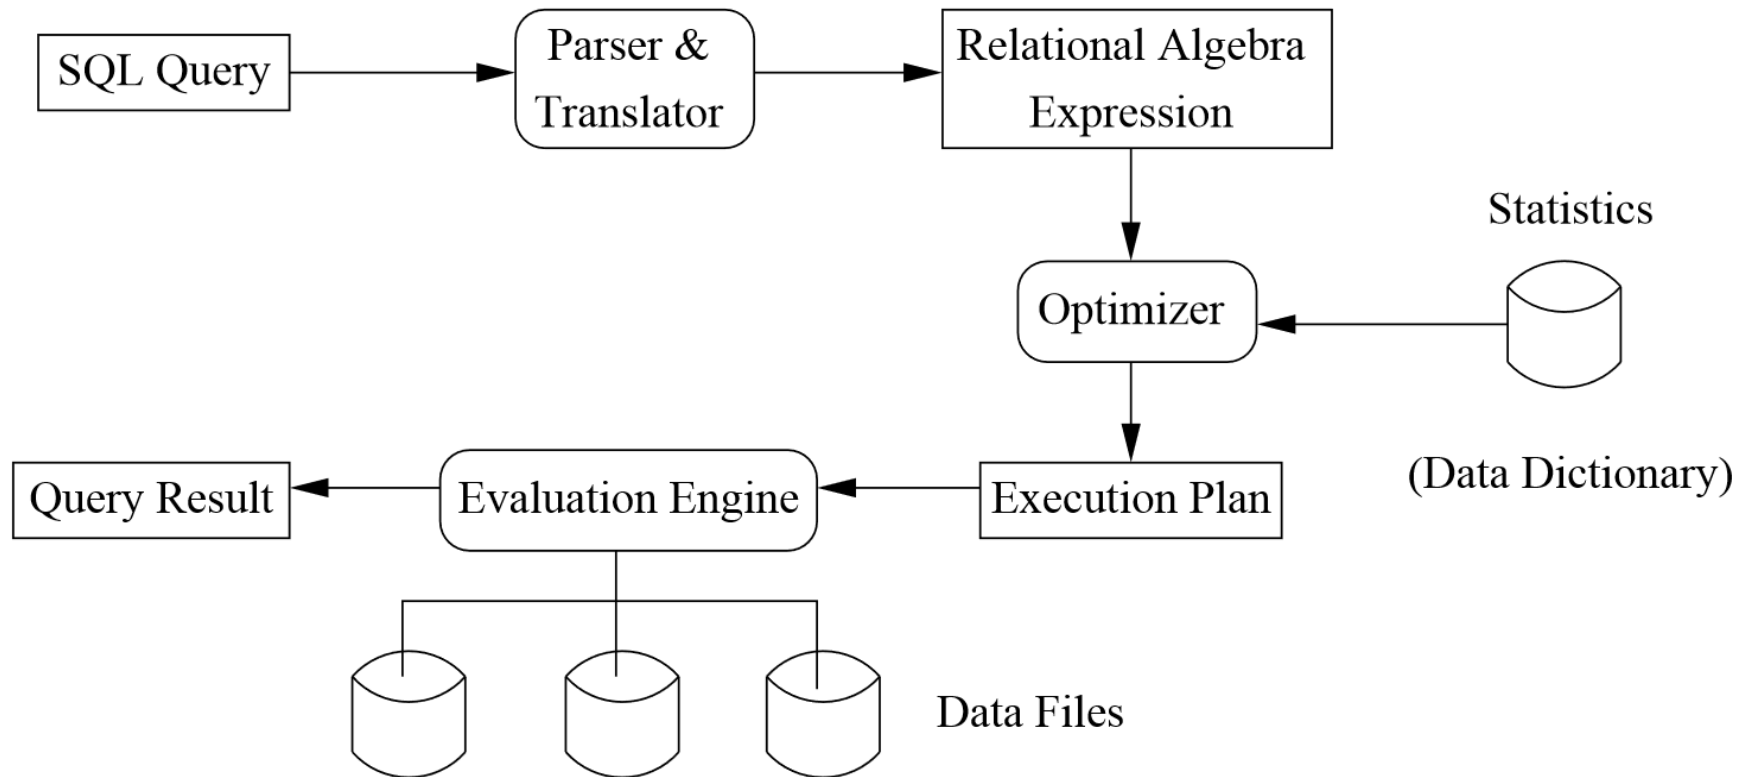
## Goal

- ▶ Minimize the *evaluation time* for a query, i.e., compute query result as fast as possible

## Cost Factors

- ▶ Disk accesses, read/write operations, [I/O, page transfer] (CPU time is typically ignored)

# Basic Steps in Processing an SQL Query



# Basic Steps in Processing an SQL Query (2)

---

## ▶ **Parsing and Translating**

- ▶ Translate the query into its internal form (parse tree).
- ▶ This is then translated into an expression of the relational algebra.
- ▶ Parser checks syntax, validates relations, attributes and access permissions

## ▶ **Optimization**

- ▶ Find the “cheapest” evaluation plan for a query

## ▶ **Evaluation**

- ▶ The query-execution engine takes a query evaluation (or query execution) plan, executes the plan, and returns the result.

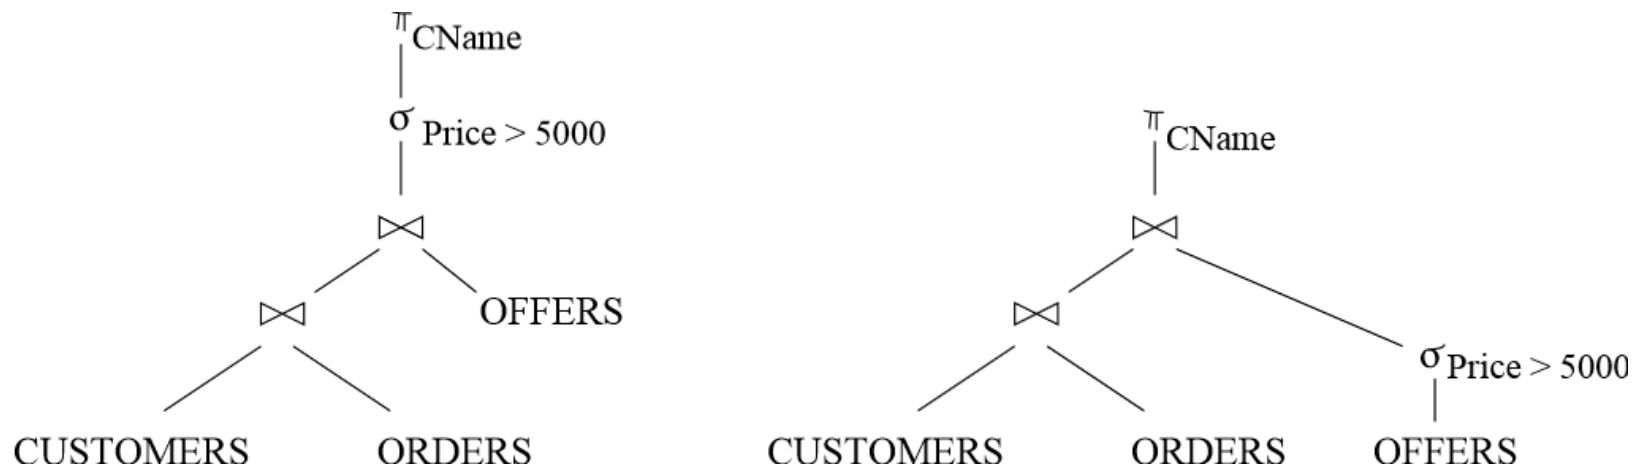
# Basic Steps in Processing an SQL Query (3)

- ▶ A relational algebra expression may have many equivalent expressions, e.g.,

$$\pi_{CName} ( \sigma_{Price > 5000} ( ( CUSTOMERS \bowtie ORDERS ) \bowtie OFFERS ) )$$

$$\pi_{CName} ( ( CUSTOMERS \bowtie ORDERS ) \bowtie ( \sigma_{Price > 5000} OFFERS ) )$$

Representation as **evaluation plan** (query tree):



Non-leaf nodes  $\equiv$  operations of relational algebra (with parameters);

Leaf nodes  $\equiv$  relations

## Basic Steps in Processing an SQL Query (4)

---

- ▶ A relational algebra expression can be evaluated in many ways.
- ▶ Annotated expression specifying detailed evaluation strategy is called **evaluation plan** (includes, e.g., whether index is used, algorithm for natural join, ...)
- ▶ Among all semantically equivalent expressions, the one **with the least costly evaluation plan** is chosen.
- ▶ Cost estimate of a plan is based on **statistical information** in the catalog (aka data dictionary).

# Catalog Information for Cost Estimation

---

## Information about relations and attributes:

- ▶  $N_R$ : number of tuples in the relation  $R$
- ▶  $B_R$ : number of blocks that contain tuples of the relation  $R$
- ▶  $S_R$ : size of a tuple of  $R$
- ▶  $F_R$ : blocking factor; number of tuples from  $R$  that fit into one block  
( $F_R = \lfloor N_R / B_R \rfloor$ )
- ▶  $V(A, R)$ : number of distinct values for attribute  $A$  in  $R$
- ▶  $SC(A, R)$ : selectivity of attribute  $A \equiv$  avg. number of tuples of  $R$  that satisfy an equality condition on  $A$ ;  
 $SC(A, R) = N_R / V(A, R)$  for non-key attributes

## Information about indexes:

- ▶  $HT_I$ : number of levels in index  $I$  ( $B^+$ -tree).
- ▶  $LB_I$ : number of blocks occupied by leaf nodes in index  $I$   
(first-level blocks).
- ▶  $Val_I$ : number of distinct values for the search key.



# Measures of Query Cost

---

- ▶ There are many possible ways to estimate cost, e.g., based on disk accesses, CPU time, or communication overhead.
- ▶ **Disk access** is the predominant cost (in terms of time);
  - ▶ relatively easy to estimate; therefore, number of block transfers from/to disk is used as measure.
  - ▶ Assumption: Each block transfer has the same cost.
- ▶ **Cost of algorithm** (e.g., for join or selection) depends on database buffer size;
  - ▶ more memory for DB buffer reduces disk accesses.
  - ▶ Thus DB buffer size is a parameter for estimating cost.
- ▶ We refer to the cost estimate of algorithm A as  $\text{cost}(A)$ . We do not consider cost of writing output to disk.

# Selection Operation

---

$\sigma_{A=a}(R)$ ,  $A$  attribute in  $R$

- ▶ **File Scan** – search algorithms that locate and retrieve records that satisfy a selection condition

- ▶ **S1** – Linear search

If  $A$  is primary key then  $\text{cost}(S1) = B_R/2$ ,  
otherwise  $\text{cost}(S1) = B_R$ .

## Selection Operation (2)

---

$\sigma_{A=a}(R)$ ,  $A$  attribute in  $R$

- ▶ **S2** – Binary search,  
i.e., the file ordered based on attribute  $A$  (primary index)

$$\text{cost}(S2) = \lceil \log_2(B_R) \rceil + \left\lceil \frac{SC(A, R)}{F_R} \right\rceil - 1$$

- ▶  $\lceil \log_2(B_R) \rceil \equiv$  cost to locate the first tuple using binary search
- ▶ Second term  $\equiv$  blocks that contain records satisfying the selection.
- ▶ If  $A$  is primary key and  $SC(A, R) = 1$ , then  $\text{cost}(S2) = \lceil \log_2(B_R) \rceil$ .
- ▶ In case of a uniform distribution of attribute values for  $A$ ,  
selection  $\sigma_{A=a}(R)$ , retrieves  $SC(A, R) = N_R/V(A, R)$  tuples.

# Selection Operation: Example

---

- ▶  $F_{\text{Employee}} = 10$ ;  
 $V(\text{Deptno}, \text{Employee}) = 50$  (different departments)
- ▶  $N_{\text{Employee}} = 10,000$  (Relation Employee has 10,000 tuples)
- ▶ Assume selection  $\sigma_{\text{Deptno}=20}(\text{Employee})$  and Employee is sorted on search key Deptno :

$\Rightarrow$

- ▶  $10,000/50 = 200$  tuples in Employee belong to Deptno 20 (assuming an equal distribution)
- ▶  $200/10 = 20$  blocks for these tuples
- ▶ A binary search finding the first block would require  $\lceil \log_2(1,000) \rceil = 10$  block accesses (assuming that the header block of the file maintains a list of file blocks)

Total cost of binary search is 10+20 block accesses (versus 1,000 for linear search and Employee not sorted by Deptno)

## Selection Operation (3)

---

**Index scan** – search algorithms that use an index (based on a B<sup>+</sup>-tree); selection condition is on search key of index

- ▶ **S3** – Primary index I for A, A primary key, equality  $A = a$   
 $\text{cost}(S3) = HT_I + 1$  (only 1 tuple satisfies condition)
- ▶ **S4** – Primary index I on non-key A, equality  $A = a$   
 $\text{cost}(S4) = HT_I + \left\lceil \frac{SC(A,R)}{F_R} \right\rceil$
- ▶ **S5** – Non-primary (non-clustered) index on non-key A, equality  $A = a$   
 $\text{cost}(S5) = HT_I + SC(A, R)$

Worst case: each matching record resides in a different block.

## Selection Operation: Example (2)

---

- ▶ Assume primary ( $B^+$ -tree) index for attribute Deptno
- ▶  $200/10=20$  blocks accesses are required to read Employee tuples
- ▶ If  $B^+$ -tree index stores 20 pointers per (inner) node, then the  $B^+$ -tree index must have between 3 and 5 leaf nodes and the entire tree has a depth of 2  
 $\Rightarrow$  a total of 22 blocks must be read.

# Selections Involving Comparisons

---

Selections of the form  $\sigma_{A \leq v}(R)$  or  $\sigma_{A \geq v}(R)$  are implemented using a file scan or binary search, or by using either a

- ▶ **S6** – A primary index on A, or
- ▶ **S7** – A secondary index on A (in this case, typically a linear file scan may be cheaper; but this depends on the selectivity of A)

# Complex Selections

---

- ▶ General pattern:

- ▶ Conjunction –  $\sigma_{\wedge_1 \wedge \dots \wedge \wedge_n} (R)$

- ▶ Disjunction –  $\sigma_{\wedge_1 \vee \dots \vee \wedge_n} (R)$

- ▶ Negation –  $\sigma_{\neg \wedge} (R)$

- ▶ The selectivity of a condition  $\Theta_i$  is the probability that a tuple in the relation  $R$  satisfies  $\Theta_i$ . If  $s_i$  is the number of tuples in  $R$  that satisfy  $\Theta_i$ , then  $\Theta_i$ 's selectivity is estimated as  $s_i/N_R$ .



# Join Operations

---

- ▶ There are several different algorithms that can be used to implement joins (natural-, equi-, condition-join)
  - ▶ Nested-Loop Join
  - ▶ Block Nested-Loop Join
  - ▶ Index Nested-Loop Join
  - ▶ Sort-Merge Join
  - ▶ Hash-Join
- ▶ Choice of a particular algorithm is based on cost estimate
- ▶ For this, join size estimates are required and in particular cost estimates for outer-level operations in a relational algebra expression.

## Join Operations: Example

---

- ▶ Assume the query  $\text{CUSTOMERS} \bowtie \text{ORDERS}$  (with join attribute only being CName)
- ▶  $N_{\text{CUSTOMERS}} = 5,000$  tuples
- ▶  $F_{\text{CUSTOMERS}} = 20$ , i.e.,  $B_{\text{CUSTOMERS}} = 5,000/20 = 250$  blocks
- ▶  $N_{\text{ORDERS}} = 10,000$  tuples
- ▶  $F_{\text{ORDERS}} = 25$ , i.e.,  $B_{\text{ORDERS}} = 400$  blocks
- ▶  $V(\text{CName}, \text{ORDERS}) = 2,500$ , meaning that in this relation, on average, each customer has four orders
- ▶ Also assume that CName in ORDERS is a foreign key on CUSTOMERS

# Estimating the Size of Joins

---

- ▶ The Cartesian product  $R \times S$  results in  $N_R * N_S$  tuples; each tuple requires  $S_R + S_S$  bytes.
- ▶ If  $\text{schema}(R) \cap \text{schema}(S) = \text{primary key for } R$ , then a tuple of  $S$  will match with at most one tuple from  $R$ . Therefore, the number of tuples in  $R \bowtie S$  is not greater than  $N_S$
- ▶ If  $\text{schema}(R) \cap \text{schema}(S) = \text{foreign key in } S$  referencing  $R$ , then the number of tuples in  $R \bowtie S$  is exactly  $N_S$
- ▶ Other cases are symmetric.

## Estimating the Size of Joins (2)

---

- ▶ In the example query CUSTOMERS  $\bowtie$  ORDERS, CName in ORDERS is a foreign key of CUSTOMERS; the result thus has exactly  $N_{\text{ORDERS}} = 10,000$  tuples
- ▶ If  $\text{schema}(R) \cap \text{schema}(S) = \{ A \}$  is not a key for R or S; assume that every tuple in R produces tuples in  $R \bowtie S$ . Then the number of tuples in  $R \bowtie S$  is estimated to be:

$$\frac{N_R * N_S}{V(A, S)}$$

- ▶ If the reverse is true, the estimate is

$$\frac{N_R * N_S}{V(A, R)}$$

and the lower of the two estimates is probably the more accurate one.

## Estimating the Size of Joins (3)

---

- ▶ Size estimates for CUSTOMERS  $\bowtie$  ORDERS without using information about foreign keys:
  - ▶  $V(\text{CName}, \text{CUSTOMERS}) = 5,000$ , and  $V(\text{CName}, \text{ORDERS}) = 2,500$
  - ▶ The two estimates are  $5,000 * 10,000 / 2,500 = 20,000$  and  $5,000 * 10,000 / 5,000 = 10,000$ .
- ▶ We choose the lower estimate, which, in this case, is the same as our earlier computation using foreign key information.

# Block Nested-Loop Join

- ▶ Evaluate the condition  $\text{join } R \bowtie_C S$

```
for each block  $B_R$  of  $R$  do begin  
  for each block  $B_S$  of  $S$  do begin  
    for each tuple  $t_R$  in  $B_R$  do  
      for each tuple  $t_S$  in  $B_S$  do  
        check whether pair  $(t_R, t_S)$   
        satisfies join condition  
        if they do, add  $t_R \circ t_S$  to the result  
    end end end end
```

- ▶  $R$  is called the *outer* and  $S$  the *inner* relation of the join.
- ▶ Requires no indexes and can be used with any kind of join condition.

## Block Nested-Loop Join (2)

---

- ▶ Worst case: db buffer can only hold one block of each relation  
 $\Rightarrow B_R + B_R * B_S$  disk accesses.
- ▶ Best case: both relations fit into db buffer  
 $\Rightarrow B_R + B_S$  disk accesses.
- ▶ If smaller relation completely fits into db buffer, use that as inner relation. Reduces the cost estimate to  $B_R + B_S$  disk accesses.
- ▶ Some improvements of block nested-loop algorithm
  - ▶ If equi-join attribute is the key on inner relation, stop inner loop with first match
  - ▶ Use  $M - 2$  disk blocks as blocking unit for outer relation, where  $M$  = db buffer size in blocks; use remaining two blocks to buffer inner relation and output.  
Reduces number of scans of inner relation greatly.
  - ▶ Scan inner loop forward and backward alternately, to make use of blocks remaining in buffer (with LRU replacement strategy)
  - ▶ Use index on inner relation, if available.

# Index Nested-Loop Join

---

- ▶ If an index is available on the inner loop's join attribute and join is an equi-join or natural join, more efficient index lookups can replace file scans.
- ▶ It is even possible (reasonable) to construct index just to compute a join.
- ▶ For each tuple  $t_R$  in the outer relation  $R$ , use the index to lookup tuples in  $S$  that satisfy join condition with  $t_R$
- ▶ Worst case: db buffer has space for only one page of  $R$  and one page of the index associated with  $S$ :
  - ▶  $B_R$  disk accesses to read  $R$ , and for each tuple in  $R$ , perform index lookup on  $S$ .
  - ▶ Cost of the join:  $B_R + N_R * c$ , where  $c$  is the cost of a single selection on  $S$  using the join condition.
- ▶ If indexes are available on both  $R$  and  $S$ , use the one with the fewer tuples as the outer relation.



# Index Nested-Loop Join (2)


---

## ► **Example:**

- Compute  $\text{CUSTOMERS} \bowtie \text{ORDERS}$ , with  $\text{CUSTOMERS}$  as the outer relation.
- Let  $\text{ORDERS}$  have a primary  $B^+$ -tree index on the join-attribute  $CName$ , which contains 20 entries per index node
- Since  $\text{ORDERS}$  has 10,000 tuples, the height of the tree is 4, and one more access is needed to find the actual data records (based on tuple identifier).
- Since  $N_{\text{CUSTOMERS}}$  is 5,000, the total cost is  $250 + 5000 * 5 = 25,250$  disk accesses.
- This cost is lower than the 100,250 accesses needed for a block nested-loop join.

# Sort-Merge Join

- ▶ Basic idea: first sort both relations on join attribute (if not already sorted this way)
- ▶ Join steps are similar to the merge stage in the external sort-merge algorithm (discussed later)
- ▶ Every pair with same value on join attribute must be matched.



- ▶ Each tuple needs to be read only once. As a result, each block is read only once. Thus, the number of block accesses is  $B_R + B_S$  plus the cost of sorting, if relations are unsorted.
- ▶ Can be used only for equi-join and natural join

## Sort-Merge Join (2)

---

- ▶ If one relation is sorted and the other has a secondary  $B^+$ -tree index on the join attribute, a *hybrid merge-join* is possible. The sorted relation is merged with the leaf node entries of the  $B^+$ -tree.
- ▶ The result is sorted on the addresses (rids) of the unsorted relation's tuples, and then the addresses can be replaced by the actual tuples efficiently.

# Hash-Join

---

- ▶ Only applicable in case of equi-join or natural join
  - ▶ A hash function is used to partition tuples of both relations into sets that have the same hash value on the join attribute
1. Partitioning Phase:  $2 * (B_R + B_S)$  block accesses
  2. Matching Phase:  $B_R + B_S$  block accesses (under the assumption that one partition of each relation fits into the database buffer)

# Cost Estimates for other Operations

---

## Sorting:

- ▶ If whole relation fits into db buffer → quick-sort
- ▶ Or, build index on the relation, and use index to read relation in sorted order.
- ▶ Relation that does not fit into db buffer → external sort-merge
  1. Phase: Create runs by sorting portions of the relation in db buffer
  2. Phase: Read runs from disk and merge runs in sort order

## Cost Estimates for other Operations (2)

---

### **Duplicate Elimination:**

- ▶ **Sorting:** remove all but one copy of tuples having identical value(s) on projection attribute(s)
- ▶ **Hashing:**
  - ▶ partition relation using hash function on projection;
  - ▶ then read partitions into buffer and create in-memory hash index;
  - ▶ tuple is only inserted into index if not already present

# Cost Estimates for other Operations (3)

---

## Set Operations:

- ▶ Sorting or hashing
- ▶ Hashing:
  - ▶ Partition both relations using the same hash function;
  - ▶ use in-memory index for partitions  $R_i$
  - ▶  $R \cup S$ : if tuple in  $R_i$  or in  $S_i$ , add tuple to result
  - ▶  $\cap$ : if tuple in  $R_i$  and in  $S_i$ , ...
  - ▶  $\neg$ : if tuple in  $R_i$  and not in  $S_i$ , ...

## Cost Estimates for other Operations (4)

---


### **Grouping:**

- ▶ Sorting or hashing.
- ▶ Hashing: while groups (partitions) are built, compute partial aggregate values (for group attribute  $A$ ,  $V(A,R)$  tuples to store values)



# Evaluation of Expressions

- ▶ The basic concept is **materialization**: Evaluate one operation at a time, starting at the lowest level. Use intermediate results materialized in temporary relations to evaluate next level operation(s).



- ▶ First compute and store  $\sigma_{Price > 5000}(OFFERS)$ ; then compute and store join of  $CUSTOMERS$  and  $ORDERS$ ; finally, join the two materialized relations and project on to  $CName$ .

## Evaluation of Expressions (2)

---

- ▶ **Pipelining**: evaluate several operations simultaneously, and pass the result (tuple- or block-wise) on to the next operation.
- ▶ In the example above, once a tuple from OFFERS satisfying selection condition has been found, pass it on to the join. Similarly, don't store result of (final) join, but pass tuples directly to projection.
- ▶ Much cheaper than materialization, because temporary relations are not generated and stored on disk.

## Evaluation of Expressions (3)

---

- ▶ Pipelining is not always possible, e.g., for all operations that include sorting (**blocking operation**).
- ▶ Pipelining can be executed in either **demand driven** or **producer driven** fashion.

# Transformation of Relational Expressions

---

- ▶ Generating a query-evaluation plan for an expression of the relational algebra involves two steps:
  1. generate logically equivalent expressions
  2. annotate these evaluation plans by specific algorithms and access structures to get alternative query plans
- ▶ Use **equivalence rules** to transform a relational algebra
- ▶ expression into an equivalent one.
- ▶ Based on estimated cost, the most cost-effective annotated plan is selected for evaluation. The process is called **cost-based query optimization**.

# Equivalence of Expressions

---

- ▶ Result relations generated by two equivalent relational algebra expressions have the same set of attributes and contain the same set of tuples, although their attributes may be ordered differently.

# Equivalence Rules (for expr. $E$ , $E_1$ , $E_2$ , cond. $F_i$ )

Applying distribution and commutativity of relational algebra operations

- 1  $\sigma_{F_1}(\sigma_{F_2}(E)) \equiv \sigma_{F_1 \wedge F_2}(E)$
- 2  $\sigma_F(E_1 \ [\cup, \cap, -] \ E_2) \equiv \sigma_F(E_1) \ [\cup, \cap, -] \ \sigma_F(E_2)$
- 3  $\sigma_F(E_1 \times E_2) \equiv \sigma_{F_0}(\sigma_{F_1}(E_1) \times \sigma_{F_2}(E_2)); F \equiv F_0 \wedge F_1 \wedge F_2, F_i$   
contains only attributes of  $E_i, i = 1, 2$ .
- 4  $\sigma_{A=B}(E_1 \times E_2) \equiv E_1 \bowtie_{A=B} E_2$
- 5  $\pi_A(E_1 \ [\cup, \cap, -] \ E_2) \equiv \pi_A(E_1) \ [\cup, \cap, -] \ \pi_A(E_2)$
- 6  $\pi_A(E_1 \times E_2) \equiv \pi_{A_1}(E_1) \times \pi_{A_2}(E_2),$  with  $A_i = A \cap \{\text{attributes in } E_i\}, i = 1, 2$ .
- 7  $E_1 \ [\cup, \cap] \ E_2 \equiv E_2 \ [\cup, \cap] \ E_1$   
 $(E_1 \cup E_2) \cup E_3 \equiv E_1 \cup (E_2 \cup E_3)$  (the analogous holds for  $\cap$ )
- 8  $E_1 \times E_2 \equiv \pi_{A_1, A_2}(E_2 \times E_1)$   
 $(E_1 \times E_2) \times E_3 \equiv E_1 \times (E_2 \times E_3)$   
 $(E_1 \times E_2) \times E_3 \equiv (E_1 \times E_3) \times E_2$
- 9  $E_1 \bowtie E_2 \equiv E_2 \bowtie E_1 \quad (E_1 \bowtie E_2) \bowtie E_3 \equiv E_1 \bowtie (E_2 \bowtie E_3)$

# Examples

---

## ► Selection:

- Find the name of all customers who have ordered a product for more than \$5,000 from a supplier located in Ilmenau.

$$\pi_{CName} (\sigma_{SAddress \text{ like } \%Ilmenau\% \wedge Price > 5000} (CUSTOMERS \bowtie (ORDERS \bowtie (OFFERS \bowtie SUPPLIERS))))$$

- Perform selection as early as possible  
(but take existing indexes on relations into account)

$$\pi_{CName} (CUSTOMERS \bowtie (ORDERS \bowtie (\sigma_{Price > 5000} (OFFERS) \bowtie (\sigma_{SAddress \text{ like } \%Ilmenau\%} (SUPPLIERS)))))$$

## Examples (2)

---

### ► Projection:

- Find the name and account of all customers who have ordered a product 'CD-ROM'

$$\pi_{\text{CName, account}} ( \text{CUSTOMERS} \bowtie \sigma_{\text{Prodname} = \text{'CD-ROM'}} ( \text{ORDERS} ) )$$

- Reduce the size of argument relation in join

$$\pi_{\text{CName, account}} ( \text{CUSTOMERS} \bowtie \pi_{\text{CName}} ( \sigma_{\text{Prodname} = \text{'CD-ROM'}} ( \text{ORDERS} ) ) )$$

Projection should not be shifted before selections, because minimizing the number of tuples in general leads to more efficient plans than reducing the size of tuples.



# Join Ordering

- ▶ For relations  $R_1, R_2, R_3$  applies  $(R_1 \bowtie R_2) \bowtie R_3 \equiv R_1 \bowtie (R_2 \bowtie R_3)$
- ▶ If  $(R_2 \bowtie R_3)$  is quite large and  $(R_1 \bowtie R_2)$  is small, we choose  $(R_1 \bowtie R_2) \bowtie R_3$  so that a smaller temporary relation is computed and materialized
- ▶ Example: List the name of all customers who have ordered a product from a supplier located in Ilmenau.

$\pi_{CName} (\sigma_{SAddress \text{ like } '%Ilmenau\%'} (SUPPLIERS \bowtie ORDERS \bowtie CUSTOMERS) )$

$ORDERS \bowtie CUSTOMERS$  is likely to be a large relation. Because it is likely that only a small fraction of suppliers are from Ilmenau, we compute the join

$\sigma_{SAddress \text{ like } '%Ilmenau\%'} (SUPPLIERS \bowtie ORDERS )$

first.

# Summary of Algebraic Optimization Rules

---

1. Perform selection as early as possible
2. Replace Cartesian Product by join whenever possible
3. Project out useless attributes early.
4. If there are several joins, perform *most restrictive* join first

# Evaluation Plan

---

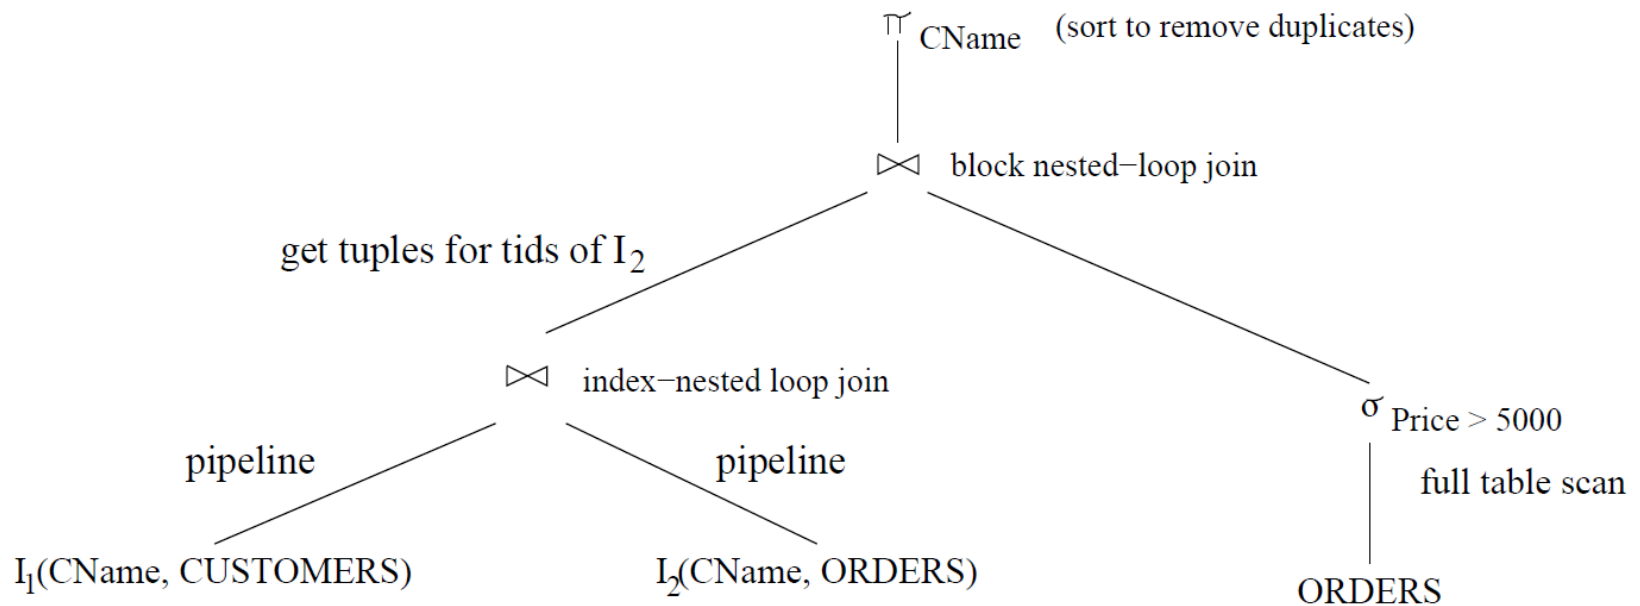
An evaluation plan for a query exactly defines

- ▶ what algorithm is used for each operation,
- ▶ which access structures are used (tables, indexes, clusters),
- ▶ and how the execution of the operations is coordinated.

# Example of Annotated Evaluation Plan

- Query: List the name of all customers who have ordered a product that costs more than \$5,000.

Assume that for both CUSTOMERS and ORDERS an index on CName exists:  $I_1(\text{CName}, \text{CUSTOMERS})$ ,  $I_2(\text{CName}, \text{ORDERS})$ .



# Choice of an Evaluation Plan

---

- ▶ Must consider interaction of evaluation techniques when choosing evaluation plan: choosing the algorithm with the least cost for each operation independently may not yield the best overall algorithm.
- ▶ Practical query optimizers incorporate elements of the following two optimization approaches:
  - ▶ **Cost-based**: enumerate all the plans and choose the best plan in a cost-based fashion.
  - ▶ **Rule-based**: Use rules (heuristics) to choose plan.
- ▶ Remarks on cost-based optimization:
  - ▶ Finding a join order for  $R_1 \bowtie R_2 \bowtie \dots \bowtie R_n$ :  
There are  $(2(n-1))! / (n-1)!$  different join orders;  
For example, for  $n=7$ , the number is 665280.  
→ use of dynamic programming techniques

## Choice of an Evaluation Plan (2)

---

- ▶ Heuristic (or rule-based) optimization transforms a given query tree by using a set of rules that typically (but not in all cases) improve execution performance:
  - ▶ Perform selection early (reduces number of tuples)
  - ▶ Perform projection early (reduces number of attributes)
  - ▶ Perform most restrictive selection and join operations before other similar operations.

# Summary

---

- ▶ Steps of query processing: translate a SQL query into an execution plan
- ▶ Cost estimation based on statistics and cost formulas
- ▶ Physical query operators: scans, join implementations, etc.
- ▶ Evaluation of query expressions
- ▶ Query optimization: search strategy for finding the best (cheapest) plan

## Example (I)

---

Customer ( Cname , Caddr , Account )  
Order ( Cname , Article , Amount )

```
SELECT Customer.Cname , Account      PROJ  
FROM Customer , Order  
WHERE Customer.Cname = Order.Cname  } SEL  
    AND Article = 'Coffee'
```

- ▶ Customer: 100 tuple ; 5 tuple / page
- ▶ Order: 10 000 tuple ; 10 tuple / page
- ▶ 50 orders related to coffee
- ▶ tuple (Customer.Cname , Account) 50 tuple / page  
⇒ result of SELECT ... → 1 page
- ▶ tuple (Customer × Order) 3 rows / page
- ▶ Buffer: size = 1 (for each relation)